



WEBASSEMBLY

What is WebAssembly

WebAssembly (abbreviated Wasm) is a binary instruction format for a stack-based virtual machine. Wasm is designed as a portable compilation target for programming languages, enabling deployment on the web for client and server applications.

What is WebAssembly? (cont.)

- Stack-based virtual machine (similar to the JVM or the CIL of .NET)
 - Non object-oriented (only supports simple datatypes)
 - Linear memory - memory is just an `ArrayBuffer` (no GC)
- Compilation target of compiled programming languages
- Sandboxed execution environment
 - Host interaction via imports/exports
 - security checks applied
- W3C standard (1.0)

How to run?

- Built-in browser support: <https://caniuse.com/?search=wasm>

Chrome	Edge *	Safari	Firefox	Opera	IE	Chrome for Android	Safari on iOS *	Samsung Internet	Opera Mini *	Opera Mobile *	UC Browser for Android	Android Browser *	Firefox for Android	QQ Browser	Baidu Browser	KaiOS Browser
			2-46													
4-50	12-14		¹ 47-51	10-37												
² 51-56	³ 15	3.1-10.1	⁴ 52	² 38-43			3.2-10.3	4-6.4								
57-124	16-124	11-17.4	53-125	44-108	6-10		11-17.4	7.2-23		12-12.1		2.1-4.4.4				¹ 2.5
125	125	17.5	126	109	11	124	17.5	24	all	80	15.5	124	125	14.9	13.52	3.1
126-128		17.6-TP	127-129				17.6									

- Built-in *node.js* support
- Dedicated / stand-alone interpreters (see later)

Which languages can compile to WASM?

- C/C++ `emscripten`
- Rust *built-in*
- Go *built-in*, `TinyGo`
- C# `Blazor`, `Uno` / F# *community*
- Kotlin *built-in*
- AssemblyScript *native*
- Dart/Flutter *built-in*
- Swift *community*
- Zig *built-in*
- Nim (via generated C-Code)
- ...

Why WASM?

- Performance
- Re-use existing (C/C++) code bases
- Hide implementation details (better than JS obfuscation)

Use Case #1: Web Apps

- Re-use existing code bases in browser apps
- Use WASM to provide calculation intense parts of an app
 - signal processing (images, video, audio)
 - Complex calculations (i.e. graphics, 3d models)
 - Weather simulation
 - Games (of course 😊)
- In use today:
 - Jupyter Notebooks (running full Python interpreter on WASM incl. numpy, ...)
 - Photoshop/Lightroom browser apps (Adobe)
 - Figma
 - TinkerCAD
 - ...

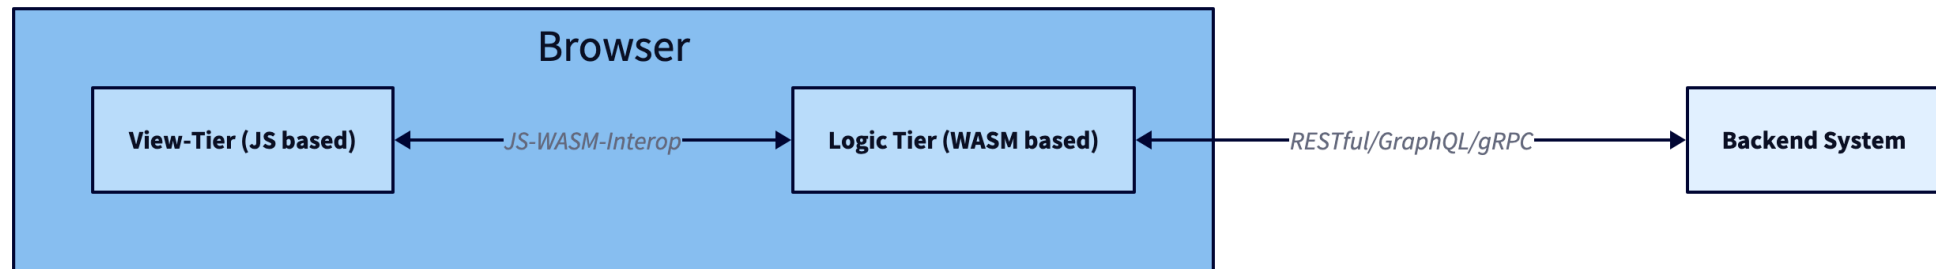
Downsides (Web Apps)

- No built-in API for DOM manipulations
 - Needs JS interaction hooks (slow)
 - Will be refined in upcoming standards draft
- "Fat" binaries
 - No GC built-in -- languages need to ship a custom GC as part of the module
 - No standard lib - languages need to package everything they need

WebAssembly 2.0 (draft) will address some of these with

- reference types
- optional GC to be used to cleanup WebAssembly memory

3-tier web app architecture



WASI - WebAssembly System Interface

- Own spec building on WebAssembly specs
- Same VM/assembly code
- Specs for types/functions that interact with the underlying OS
 - Filesystem
 - I/O streams
- Very early draft (current spec is named "preview1", "preview2" is in dev)

Use Case #2: Backend Apps/Serverless/FaaS

- Write portable backend apps/functions targetting WASM/WASI
 - Use any language that can compile to WASM/WASI
- Different VM implementations: wasmtime, WAMR, WasmEdge, wazero, Wasmer, wasmi, wasm3
- Kubernetes devs work on providing k8s for WASM (instead of Containers): FaaS

Downsides (Backend Apps)

- WASI specs early drafts
- Lots of features missing (i.e. direct network access)
- Area of active development

How does a virtual stack machine work

- operands are pushed onto a stack
- Operations pop operands off the stack and push the result back onto the stack
- Function arguments and return values are also passed via the stack

How does a virtual stack machine work (cont.)

2 + 3

```
i32.const 2  
i32.const 3  
i32.add
```

How does a virtual stack machine work (cont.)

2 + 3

```
i32.const 2  
i32.const 3  
i32.add
```

... or equivalent as *S-expressions*

```
( i32.add  
  ( i32.const 2 )  
  ( i32.const 3 )  
)
```

How does it look like?

```
(module
  ;; Define a function fib to calculate the n-th fibonacci number in a resource wasting way.
  (func $fib (export "fib") (param i64) (result i64)
    ;; Calculate param < 3
    (i64.lt_u
      (local.get 0)
      (i64.const 3)
    )

    ;; If result of comparison is true (!= 0)...
    (if (result i64)
      (then
        ;; ... then put 1 on the stack ...
        i64.const 1
      )
      (else
        ;; ... otherwise compute fib(param - 1) + fib(param - 2)
        ;; and leave the result on the stack
        (i64.add
          (call $fib
            (i64.sub
              (local.get 0)
              (i64.const 1)
            )
          )
          (call $fib
            (i64.sub
              (local.get 0)
              (i64.const 2)
            )
          )
        )
      )
    )
  )
)
```


Hello, WASI

```
(module
  ;; Import the required fd_write WASI function which will write the given io vectors to stdout
  ;; The function signature for fd_write is:
  ;; (File Descriptor, *iovs, iovec_len, *nwritten) -> Returns 0 on success, nonzero on error
  (import "wasi_snapshot_preview1" "fd_write" (func $fd_write (param i32 i32 i32 i32) (result i32)))

  ;; Reserve 1 page of memory (=64Kb) and export it
  (memory 1)
  (export "memory" (memory 0))

  ;; Write 'hello world\n' to memory at an offset of 8 bytes
  (data (i32.const 8) "hello world\n")

  (func (export "_start")
    ;; Creating a new io vector within linear memory
    (i32.store (i32.const 0) (i32.const 8)) ;; iov.iov_base - This is a pointer to the start of the 'hello world\n' string
    (i32.store (i32.const 4) (i32.const 12)) ;; iov.iov_len - The length of the 'hello world\n' string

    (call $fd_write
      (i32.const 1) ;; file_descriptor - 1 for stdout
      (i32.const 0) ;; *iovs - The pointer to the iov array, which is stored at memory location 0
      (i32.const 1) ;; iovec_len - We're printing 1 string stored in an iov - so one.
      (i32.const 20) ;; nwritten - A place in memory to store the number of bytes written
    )
    drop ;; Discard the number of bytes written from the top of the stack
  )
)
```

How fast is WASM?

Very simple CPU-intensive benchmark:

```
// Javascript
function fib (x) {
  if (x === 0) return 0;
  if (x < 2) return 1;
  return fib(x-1) + fib(x-2);
}
```

```
// C
unsigned long long fib(const unsigned long long x) {
  if (x == 0) return 0;
  if (x < 3) return 1;
  return fib(x-1) + fib(x - 2);
}
```

How fast is WASM? (cont.)

Implementation	Time for <code>fib(45)</code> [s]	Size of the binary [kB]	env
JavaScript	6.3	0.2	node v22
C (native; unoptimized)	4.6	33	macOS
C (native; optimized)	2.0	33	macOS
C (WASM; unoptimized)	9.3	$(59 + 13) = 72$	node v22
C (WASM; optimized)	2.7	$(12 + 6.3) = 18.3$	node v22
WASM (handwritten)	4.7	0.06	wasmtime

How fast is WASM? (cont.)

Implementation	Time for <code>fib(45)</code> [s]	Size of the binary [kB]	env
Go (native)	2.4	2,000	macOS
Go (WASM)	12.2	2,100	node v22 w/ custom JS
Go (WASM)	47.2	2,100	wasmtime
TinyGo (WASM)	2.9	603	node v22 w/ custom JS
TinyGo (WASM)	2.8	603	wasmtime

```
function fib (x) {  
  if (x < 2) return x;  
  let a = 0, b = 1;  
  for (let i = 2; i <= x; i++) {  
    const tmp = a + b; a = b; b = tmp;  
  }  
  return b;  
}
```

Implementation	Time for <code>fib(45)</code> [s]
JavaScript	6.3
C (native; optimized)	2.0
JavaScript (dynamic programming; bottom up)	0.05

Live Demo

Stuff to check out

- <https://webassembly.org>
- <https://developer.mozilla.org/en-US/docs/WebAssembly>
- <https://wasi.dev>
- <https://www.cncf.io/blog/2024/03/12/webassembly-on-kubernetes-from-containers-to-wasm-part-01/>